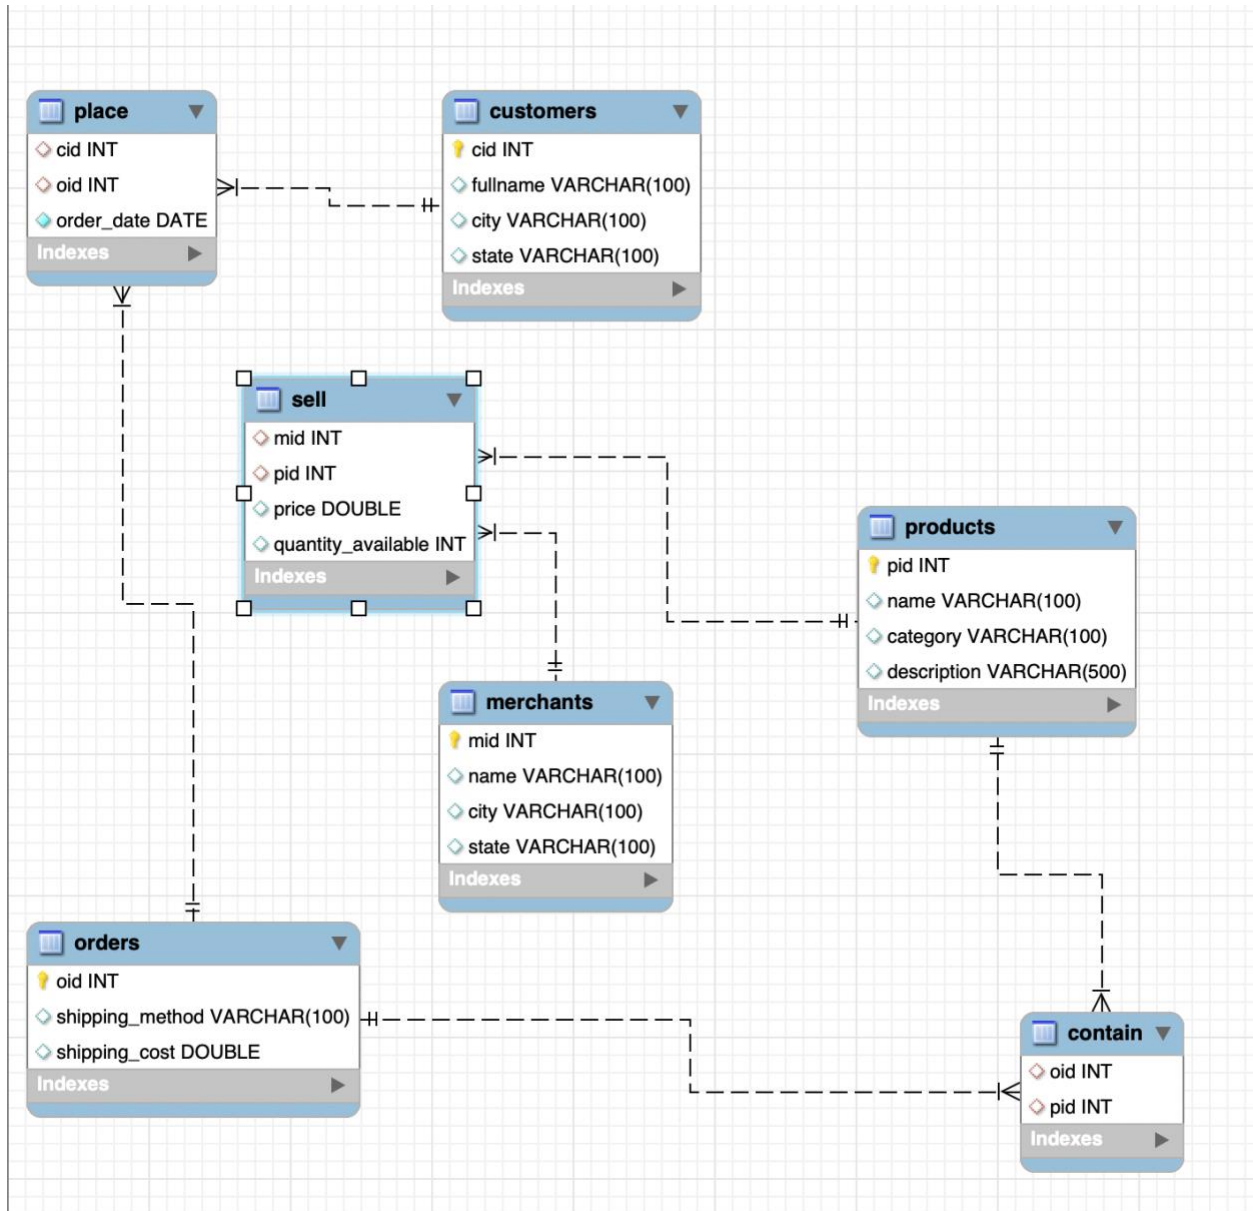


Title: DB Assignment 3

Name: Devon Reing

Date: 8 October 2024



The ER diagram for this picture illustrated above shows the relationships between the 7 tables and the attributes for each entity set. The tables sell, contain, and place are all relationships that connect the remaining entity sets which can be seen through the foreign keys. Each of the foreign keys comes from the connected entity's primary key.

```
78  -- Question 1
79  * select products.name, merchants.name
80  from products
81      join sell on products.pid = sell.pid
82      join merchants on merchants.mid = sell.mid
83  where sell.quantity_available = 0;
84
85  Question 2
```

100% 35:83 1 error found

Result Grid Filter Rows: Search Export:

| | name | name | |
|--|------------------|--------|--|
| | Router | Acer | |
| | Network Card | Acer | |
| | Printer | Apple | |
| | Router | Apple | |
| | Router | HP | |
| | Super Drive | HP | |
| | Laptop | HP | |
| | Router | Dell | |
| | Ethernet Adapter | Lenovo | |

The first query aims to list every product that is no longer available. This is done through checking when the quantity available is 0 in the sell table for each product. To get each product along with who sells it an inner join is done on products, sell, and merchants through the id fields for each respective table. The attributes shown in the output is provided by the select statement, which provides both the merchant name and the product name.

```
85 -- Question 2
86 -- query used to check
87 select *
88 from products
89     left join sell
90     on products.pid = sell.pid;
91 -- query that returns the intended result from question
92 select products.name, products.description
93 from products
94     left join sell
95     on products.pid = sell.pid
96 where sell.pid is null; -- check if null because that will list products that have no merchants selling them
97
```

100% 24:96 1 error found

Result Grid Filter Rows: Search Export:

| | name | description |
|--|-------------|--------------------|
| | Super Drive | External CD/DVD/RW |
| | Super Drive | Internal CD/DVD/RW |

The second query aims to list the products in the products table that are not sold by any merchant. To achieve this, a left join is done on products with sell using the product id attributes in each table. With a left join, every entry from the products (left side) will be kept in the new combined table. If an entry with the same product id is not present in the sell table, the remaining columns from the sell table will be filled with null values. Using the where statement, the products to appear in the output are filtered by checking for a null value in the sell column for product id. To format the output and avoid all columns from appearing in the output table, the select statement limits the columns that appear in the output table to the products name and description.

```
98      -- Question 3
99      select count(
100      ((select DISTINCT customers.cid
101      from customers
102          join place on customers.cid = place.cid
103          join contain on place.oid = contain.oid
104          join products on products.pid = contain.pid
105      where products.description LIKE '%SATA%')
106      except
107      (select DISTINCT customers.cid
108      from customers
109          join place on customers.cid = place.cid
110          join contain on place.oid = contain.oid
111          join products on products.pid = contain.pid
112      where products.name = 'Router')));
113
```

100% 35:112 1 error found

Result Grid Filter Rows: Search Export:

| |
|---|
| count(((select DISTINCT customers.cid |
| 0 |

The third query aims to find the number of customers that bought SATA drive products but not any routers. To achieve this, two separate subqueries are executed and joined together through an except statement to only show the values that appear in the first subquery and not the second. The first subquery selects the distinct customers that have purchased a drive with SATA in the description. To make this possible, customers, place, contain, and products are combined using an inner join on the id values present in the tables. The distinct keyword ensures a customer is not included twice in the list if they have bought multiple products that satisfy the criteria. The second subquery follows a similar process and combines the same tables in the same way. Then each distinct customer id that has purchased a product with the name router is selected. The except statement removes any values returned in the second subquery from the results returned from the first subquery. Finally the main query takes the returned customer ids and counts the number of them that are present to return that in the output table.

```

120  -- Question 4
121  • START TRANSACTION; -- allows for resetting the prices back to full price incase query was wrong
122  • UPDATE sell
123      join merchants on merchants.mid = sell.mid
124      join products on products.pid = sell.pid
125      SET sell.price = (sell.price * .8)
126      WHERE merchants.name = 'HP' AND products.category = 'Networking';
127
128  • SELECT sell.mid, sell.pid, merchants.name, products.category, sell.price
129      from sell
130      join merchants on sell.mid = merchants.mid
131      join products on products.pid = sell.pid
132      where merchants.name = 'HP' and products.category = 'Networking';
133  • ROLLBACK; -- incase update goes wrong
134
100% 67:132 1 error found

```

Result Grid

| mid | pid | name | category | price |
|-----|-----|------|------------|--------------------|
| 3 | 8 | HP | Networking | 827.5680000000001 |
| 3 | 10 | HP | Networking | 923.7440000000001 |
| 3 | 12 | HP | Networking | 276.008 |
| 3 | 13 | HP | Networking | 209.76 |
| 3 | 16 | HP | Networking | 1008.3600000000001 |
| 3 | 18 | HP | Networking | 164.448 |
| 3 | 19 | HP | Networking | 1179.896 |
| 3 | 20 | HP | Networking | 441.616 |
| 3 | 23 | HP | Networking | 80.76 |
| 3 | 28 | HP | Networking | 943.2080000000001 |

The fourth query aims to provide a 20% discount for every HP networking product. To achieve this, an update query is executed where the price in the sell table is multiplied by .8 to set the new price to be 80% of the original price. In order to ensure only the HP networking products are affected by this, sell is joined with merchants and products using an inner join on the id attributes. This allows the where statement to access the merchant name attribute and product category to make sure they are both equal to HP and networking respectively. To show these changes in the output table, the same joins and where statement are executed again in a select statement where the merchant id, product id, merchant name, product category, and new product price are selected to be columns in the output table. Finally the start transaction and rollback keywords surround the statements. The start transaction keyword makes it so that the original table data is saved somewhere incase an update goes wrong. If the table data needs to be reset to the original data, the rollback statement can be executed.

```

135  -- Question 5
136  • SELECT customers.fullname, products.name, sell.price
137      from customers
138          join place on customers.cid = place.cid
139          join contain on place.oid = contain.oid
140          join products on contain.pid = products.pid
141          join sell on products.pid = sell.pid
142          join merchants on sell.mid = merchants.mid
143      where customers.fullname = 'Uriel Whitney' and merchants.name = 'Acer';
144
100% 72:143 1 error found

```

| fullname | name | price |
|---------------|--------------|---------|
| Uriel Whitney | Monitor | 1435.38 |
| Uriel Whitney | Router | 521.07 |
| Uriel Whitney | Router | 1256.57 |
| Uriel Whitney | Monitor | 1103.47 |
| Uriel Whitney | Super Drive | 356.13 |
| Uriel Whitney | Printer | 1345.37 |
| Uriel Whitney | Super Drive | 671.75 |
| Uriel Whitney | Super Drive | 1135.3 |
| Uriel Whitney | Super Drive | 356.13 |
| Uriel Whitney | Super Drive | 1015.95 |
| Uriel Whitney | Network C... | 405.4 |
| Uriel Whitney | Hard Drive | 836.99 |
| Uriel Whitney | Super Drive | 1124.26 |
| Uriel Whitney | Network C... | 609.2 |
| Uriel Whitney | Printer | 1345.37 |
| Uriel Whitney | Network C... | 405.4 |
| Uriel Whitney | Super Drive | 671.75 |
| Uriel Whitney | Super Drive | 1135.3 |
| Uriel Whitney | Router | 945.51 |
| Uriel Whitney | Hard Drive | 333.71 |
| Uriel Whitney | Laptop | 247.96 |
| Uriel Whitney | Router | 394.04 |
| Uriel Whitney | Laptop | 33.5 |
| Uriel Whitney | Super Drive | 1015.95 |
| Uriel Whitney | Super Drive | 671.75 |
| Uriel Whitney | Router | 521.07 |

Result 8

The fifth query aims to list every retrieve the orders Uriel Whitney placed from Acer. To achieve this, customers, place, contain, products, sell, and merchants are joined using an inner join on the id attributes. Using this new combined table, the customers full name, product name and product price is selected to be in the output table. To narrow down the data in the output table, the where statement selects on the rows that have the customer full name listed as Uriel Whitney and the merchant name listed as Acer.

147

-- Question 6

148

select merchants.name, YEAR(place.order_date), sum(sell.price)

149

from merchants

150

join sell on merchants.mid = sell.mid

151

join contain on contain.pid = sell.pid

152

join place on place.oid = contain.oid

153

group by merchants.name, YEAR(place.order_date)

154

order by YEAR(place.order_date) desc;

155

100%

38:154

1 error found

Result Grid

Filter Rows:

Search

Export:

| | name | YEAR(place.order_dat... | sum(sell.price) | |
|--|--------|-------------------------|--------------------|--|
| | Acer | 2020 | 182311.14999999994 | |
| | Apple | 2020 | 216461.06000000006 | |
| | Dell | 2020 | 208063.07999999987 | |
| | Lenovo | 2020 | 214154.25000000002 | |
| | HP | 2020 | 164084.18200000003 | |
| | Acer | 2019 | 208815.79999999993 | |
| | Apple | 2019 | 231573.17000000007 | |
| | HP | 2019 | 156175.70400000003 | |
| | Dell | 2019 | 221391.82999999975 | |
| | Lenovo | 2019 | 232610.80000000001 | |
| | Acer | 2018 | 262059.28999999998 | |
| | Apple | 2018 | 300413.22999999986 | |
| | Dell | 2018 | 315004.82 | |
| | Lenovo | 2018 | 324291.59000000067 | |
| | HP | 2018 | 202028.69200000002 | |
| | Acer | 2017 | 176722.76999999987 | |
| | Apple | 2017 | 179560.78000000003 | |
| | Dell | 2017 | 182288.60999999996 | |
| | Lenovo | 2017 | 197980.33000000013 | |
| | HP | 2017 | 124796.61799999999 | |
| | Apple | 2016 | 64748.45999999995 | |
| | HP | 2016 | 52167.494000000006 | |
| | Dell | 2016 | 71462.86999999998 | |
| | Lenovo | 2016 | 70131.56999999998 | |
| | Acer | 2016 | 60291.140000000014 | |
| | Acer | 2011 | 152986.29999999993 | |
| | Apple | 2011 | 166822.90999999995 | |
| | HP | 2011 | 128208.28 | |
| | Dell | 2011 | 181730.34999999998 | |
| | Lenovo | 2011 | 184939.41000000006 | |

The sixth query aims to list the total annual sales for each company. This is achieved by selecting the merchant name, year from the order date attribute from the order table, and summing the price of the products sold. In order to be able to get all this information, an inner join using the id

attribute is done on the merchants, sell, contain, and place tables. To ensure the product prices are added correctly according to each company and year, a group by statement specifies that the sum should be grouped according to merchant name and the year in the order date attribute. Finally, to make the output table easily readable, the output table is ordered by the year in descending order to show the most recent results at the top.

```
156 -- Question 7
157 • select merchants.name, YEAR(place.order_date), sum(sell.price) as total_sales
158 from merchants
159     join sell on merchants.mid = sell.mid
160     join contain on contain.pid = sell.pid
161     join place on place.oid = contain.oid
162 group by merchants.name, YEAR(place.order_date)
163 having total_sales >= all
164     (select sum(sell.price)
165      from merchants
166          join sell on merchants.mid = sell.mid
167          join contain on contain.pid = sell.pid
168          join place on place.oid = contain.oid
169          group by merchants.name, YEAR(place.order_date));
170
```

100% 51:169 1 error found

Result Grid Filter Rows: Search Export:

| name | YEAR(place.order_dat... | total_sales |
|--------|-------------------------|--------------------|
| Lenovo | 2018 | 324291.59000000067 |

The seventh query aims to determine which company had the highest annual revenue and in what year that occurred. The same query besides the order by statement is used to begin this query. To distinguish this query from the sixth one a having statement is included. The having statement selects the result from the sixth query that has the highest total sales through a subquery. The subquery determines the sum of the prices of the products sold for each company in each year. All these results are then compared through the having statement and only the results that are greater than or equal to all the other entries are added to the output table. This is similar to using

max() but allows for multiple results to be returned if two or more results held the same maximum value.

```
168  -- Question 8
169  •  select orders.shipping_method, avg(orders.shipping_cost)
170      from orders
171      group by orders.shipping_method
172      having avg(orders.shipping_cost) <= all
173         (select avg(orders.shipping_cost)
174          from orders
175          group by orders.shipping_method);
```

100% 38:175 1 error found

Result Grid Filter Rows: Search Export:

| shipping_meth... | avg(orders.shipping_c... |
|------------------|--------------------------|
| USPS | 7.455760869565214 |

The eighth query aims to determine the average cheapest shipping method. This is done by selecting the shipping method and performing the average of the shipping cost for that shipping method. The average is grouped by shipping method to ensure that the average is computed for each shipping method and not the average of all shipping methods combined. Since the shipping method and cost is only available through the orders table and the rows in the orders table only exist if an order is placed, no joins need to be performed to make sure only shipping methods that have actually been used are included in the calculation. To output only the cheapest method, a having statement is executed where the average shipping cost grouped by shipping method is again computed in a subquery. Then the results from that are compared to one another and the having statement makes it so that only the minimum value is outputted to the output table. If multiple equal minimum values exist, all of those will be provided in the output table.

```

180  -- Question 9
181  with category_sales as (
182      select merchants.mid, merchants.name as merchant_name, products.category, sum(sell.price * sell.quantity_available) as totals
183      from merchants
184      join sell on merchants.mid = sell.mid
185      join products on sell.pid = products.pid
186      join contain on contain.pid = products.pid
187      group by merchants.mid, products.category
188  )
189  select cs.merchant_name, cs.category, cs.totals
190  from category_sales cs
191  join (
192      select mid, max(totals) as max
193      from category_sales
194      group by mid
195  ) max_sales on cs.mid = max_sales.mid AND cs.totals = max_sales.max;

```

100% 70:195 1 error found

Result Grid Filter Rows: Search Export:

| merchant_name | category | totals |
|---------------|------------|--------------------|
| Acer | Peripheral | 5281119.9899999695 |
| Apple | Peripheral | 3938546.6099999878 |
| HP | Peripheral | 3055029.310000003 |
| Dell | Peripheral | 6338444.070000001 |
| Lenovo | Peripheral | 5336522.410000023 |

The ninth query aims to determine what the best sold category for each company is based on the amount of money that category has produced for the company in sales. To accomplish this, first a CTE table is created as category sales. This table selects the merchant id and name, product category, and the sum of the price of products sold. The merchant name receives an alias of merchant_name and the sum receives an alias of totals. This select statement is possible through the inner join between merchants, sell, products, and contain on the id attributes. The sum is calculated for each category for each merchant through the group by statement, and the results of this output table are saved as the category sales table to be used in the main query. The main query selects the merchant name, category, and totals from the columns of the CTE table. It is then joined with a table formed through a subquery. The subquery selects the merchant id and the maximum of the totals for each merchant. The tables from the CTE and subquery are then joined using an inner join where the merchant ids are the same and the sum from the CTE table and the max total returned from the subquery are both the same. This ensures that each company is listed and only their highest selling category is included in the output table.

```

-- Question 10
-- max and min customer at each company
209
210 with customer_spending as (
211     select merchants.mid, merchants.name, customers.cid, customers.fullname, SUM(sell.price) as total_spent
212     from merchants
213         join sell on merchants.mid = sell.mid
214         join products on sell.pid = products.pid
215         join contain on products.pid = contain.pid
216         join place on contain.oid = place.oid
217         join customers on place.cid = customers.cid
218     group by merchants.mid, customers.cid, customers.fullname
219 )
220 select cs.name, cs.fullname, cs.total_spent
221 from customer_spending cs
222 where cs.total_spent in (
223     select MAX(total_spent)
224     from customer_spending
225     where mid = cs.mid
226     group by mid
227 )
228 OR
229 cs.total_spent in (
230     select MIN(total_spent)
231     from customer_spending
232     where mid = cs.mid
233     group by mid
234 )
235 order by cs.mid, cs.total_spent;
236
100% 33:235 1 error found

```

| | name | fullname | total_spent |
|--|--------|-------------------|--------------------|
| | Acer | Inez Long | 31901.019999999993 |
| | Acer | Dean Heath | 75230.289999999998 |
| | Apple | Inez Long | 32251.099999999998 |
| | Apple | Clementine Travis | 84551.109999999997 |
| | HP | Inez Long | 23483.798000000003 |
| | HP | Clementine Travis | 60203.294 |
| | Dell | Inez Long | 31135.740000000001 |
| | Dell | Clementine Travis | 85611.549999999999 |
| | Lenovo | Inez Long | 33948.909999999996 |
| | Lenovo | Haviva Stewart | 83030.259999999997 |

The tenth query aims to find out what customers have spent the most and least amounts for each company. The query begins with creating a CTE table named customer spending. This table selects the merchant id and name, customer id and name, and finds the sum of the prices of products sold under the alias total_spent. These selects are possible through inner joins on the id attribute between the merchants, sell, products, contain, place, and customers tables. The sum is computed individually for each company and customer through the group by statement which specifies that the sums should be grouped by company first and then by customer. After this table is created with the attributes in the select statement as the columns, the main query begins by

selecting the merchant name, customer name, and total spent columns from the CTE table. The rows returned from the select statement are then narrowed down to include only the maximum and minimum spenders through the where statement in the main query. The max is calculated first through a subquery in the where statement. This subquery selects the max entry in total spent column from the CTE table where the merchant id matches in the subquery and the CTE table. This is done for each company through the group by statement that specifies it should be grouped together by merchant id. The same process is then followed in a second subquery in the where statement for the minimum value. Both of these values are returned in the output table for each company through the or keyword in the where statement between the subqueries. Finally, the output table is ordered by the minimum then maximum value for each company, making sure that each company's results are outputted in subsequent rows together.